

Chapter 4

Finite Automata

Another way of characterizing regular languages are finite automata (FA)[18]. We will show that the languages of finite automata are exactly the regular languages. Furthermore, we will also derive a decision procedure for equivalence of regular expressions.

4.1 Definition

A finite automaton consists of

1. finite set of states Q ,
2. a starting state $s \in Q$,
3. a set of final states $F \subseteq Q$
4. and a state-transition relation δ .

We define a **run** of a word $w \in \Sigma^*$ on an automaton $A = (Q, s, F, \delta)$ as a sequence of states σ such that for every two consecutive positions $i, i + 1$ in σ there is $(\sigma_i, w_i, \sigma_{i+1}) \in \delta$. A word w is **accepted** by A in state x if and only if there exists a run σ of w on A such that $\sigma_0 = x \wedge \sigma_{|\sigma|-1} \in F$. The resulting set of accepted words is denoted by $\mathcal{L}_x(A)$. The **language** of A is exactly $\mathcal{L}_s(A)$ and is denoted $\mathcal{L}(A)$.

4.1.1 Non-Deterministic Finite Automata

Finite automata can be **non-deterministic** (NFA) in the sense that there exist multiple distinct runs for a word. This is the case if and only if δ is not functional.

Listing 4.1: Non-Deterministic Finite Automata

```
Record nfa : Type :=  
{ nfa_state := finType;
```

```

nfa_s : nfa_state ;
nfa_fin : pred nfa_state ;
nfa_step : nfa_state -> char -> pred nfa_state }.
Fixpoint nfa_accept (x: A) w :=
match w with
| [] => nfa_fin A x
| a::w => [ exists y, ( nfa_step A x a y ) && nfa_accept y w ]
end.
Definition nfa_lang := [pred w | nfa_accept (nfa_s A) w].

```

The acceptance criterion given here avoids the matter of runs. In many cases, this will help us with proofs by induction on the accepted word. However, we will need runs in some of the proofs. Due to the fact that runs are not unique on NFAs, we give a predicate that decides if a run on A is valid for a word w . We then show that the acceptance criterion given above corresponds to the mathematical definition in terms of runs.

```

Fixpoint nfa_run x (xs : seq A) (w: word) {struct xs} :=
match xs,w with
| y :: xs', a::w' => nfa_step A x a y && nfa_run y xs' w'
| [] , [] => true
| - , - => false
end.
Lemma nfa_run_accept x w:
  reflect ( exists2 xs, nfa_run x xs w & last x xs \in nfa_fin A )
    (nfa_accept x w).

```

4.1.2 Deterministic Finite Automata

For functional δ , we speak of **deterministic** finite automata (DFA). In this case, we write δ as a function in our COQ development.

Listing 4.2: Deterministic Finite Automata

```

Record dfa : Type :=
{ dfa_state :> finType;
  dfa_s : dfa_state ;
  dfa_fin : pred dfa_state ;
  dfa_step : dfa_state -> char -> dfa_state }.
Fixpoint dfa_accept x w :=
match w with
| [] => dfa_fin A x
| a::w => dfa_accept (dfa_step A x a) w
end.
Definition dfa_lang := [pred w | dfa_accept (dfa_s A) w].

```

Again, we avoid runs in our formalization of the acceptance criterion in favor of a acceptance criterion that is easier to work with in proofs. In this case, however, we can give a function that computes the unique run of a word on A . This allows us to give an alternative acceptance criterion that is closer to the mathematical definition. We also prove that both criteria are equivalent.

Fixpoint $\text{dfa_run}' (x: A) (w: \text{word}) : \text{seq } A :=$
match w **with**
 | $[:]$ $=>$ $[:]$
 | $a :: w =>$ $(\text{dfa_step } A \times a) :: \text{dfa_run}' (\text{dfa_step } A \times a) w$
end.
Lemma $\text{dfa_run_accept } x \ w: \text{last } x \ (\text{dfa_run}' \ x \ w) \ \backslash \text{in } \text{dfa_fin } A = (w \ \backslash \text{in } \text{dfa_accept } x).$

Equivalence of Automata

Definition 4.1.1. *We say that two automata are **equivalent** if and only if their languages are equal.*

Equivalence between DFA and NFA

Deterministic and non-deterministic finite automata are equally expressive. One direction is trivial since every DFA can be seen as a NFA. We prove the other direction using the powerset construction. Given NFA A , we construct an equivalent DFA A_{det} in the following way:

$$\begin{aligned} Q_{det} &:= 2^Q \\ s_{det} &:= \{s\} \\ F_{det} &:= \{P \mid P \in Q_{det} \wedge P \cap F \neq \emptyset\} \\ \delta_{det} &:= \{(P, a, \bigcup_{p \in P} \{q \mid q \in Q, (p, a, q) \in \delta\}) \mid P \in Q_{det}, a \in \Sigma\}. \\ A_{det} &:= (Q_{det}, s_{det}, F_{det}, \delta_{det}). \end{aligned}$$

The formalization of A_{det} is straight-forward. We leave the set of states $\{\text{set } A\}$ implicit.

Definition $\text{nfa_to_dfa} :=$
 $\{ \mid \text{dfa_s} := \text{set1 } (\text{nfa_s } A);$
 $\text{dfa_fin} := [\text{pred } X: \{\text{set } A\} \mid [\text{exists } x: A, (x \ \backslash \text{in } X) \ \&\& \ \text{nfa_fin } A \ x]];$
 $\text{dfa_step} := [\text{fun } X \ a => \ \backslash \text{bigcup_}(x \mid x \ \backslash \text{in } X) \ \text{finset } (\text{nfa_step } A \ x \ a)] \ \mid \}$
 $.$

Lemma 4.1.2. *For all powerset states X and for all states x with $x \in X$ we have that*

$$\mathcal{L}_x(A) \subseteq \mathcal{L}_X(A_{det}).$$

Proof. Let $w \in \mathcal{L}_x(A)$. We proof by induction on w that $w \in \mathcal{L}_X(A_{det})$.

- For $w = \varepsilon$ and $\varepsilon \in \mathcal{L}_x(A)$ we get $x \in F$ from $\varepsilon \in \mathcal{L}_x(A)$. From $x \in X$ we get $X \cap F \neq \emptyset$ and therefore $\varepsilon \in \mathcal{L}_X(A_{det})$.
- For $w = aw'$ and $aw' \in \mathcal{L}_x(A)$ we get y such that $w' \in \mathcal{L}_y(A)$ and $(x, a, y) \in \delta$. The latter gives us $y \in Y$ where Y is such that $(X, a, Y) \in \delta_{det}$. With $y \in Y$ and $w' \in \mathcal{L}_y(A)$ we get which gives us $w' \in \mathcal{L}_Y(A_{det})$ by induction hypothesis. With $(X, a, Y) \in \delta_{det}$ we get $aw' \in \mathcal{L}_X(A_{det})$.

□

Lemma 4.1.3. *For all powerset states X and all words $w \in \mathcal{L}_X(A_{det})$ there exists a state x such that*

$$x \in X \wedge w \in \mathcal{L}_x(A).$$

Proof. We do an induction on $w \in \Sigma^*$.

- For $w = \varepsilon$ and $\varepsilon \in \mathcal{L}_X(A_{det})$ we get $X \cap F \neq \emptyset$. Therefore, there exists x such that $x \in X$ and $x \in F$. Thus, we have $\varepsilon \in \mathcal{L}_x(A)$.
- For $w = aw'$ and $aw' \in \mathcal{L}_X(A_{det})$ we get Y such that $w' \in \mathcal{L}_Y(A_{det})$ and $(X, a, Y) \in \delta_{det}$. From the induction hypothesis we get y such that $y \in Y$ and $w' \in \mathcal{L}_y(A)$. From $y \in Y$ and $(X, a, Y) \in \delta_{det}$ we get x such that $x \in X$ and $(x, a, y) \in \delta$. Thus, $aw' \in \mathcal{L}_x(A)$.

□

Theorem 4.1.4. *The powerset automaton A_{det} accepts the same language as A , i.e.*

$$\mathcal{L}(A) = \mathcal{L}(A_{det}).$$

Proof. “ \subseteq ” This follows directly from lemma 4.1.2 with $x = s$ and $X = s_{det}$.

“ \supseteq ” From lemma 4.1.3 with $X = s_{det}$ we get $\mathcal{L}_{s_{det}}(A_{det}) \subseteq \mathcal{L}_s(A)$, which proves the claim. □

The formalization of this proof is straight-forward and follows the plan laid out above. The corresponding Lemmas are:

Lemma `nfa_to_dfa_complete` (x : A) w (X : `nfa_to_dfa`):

`x \in X \rightarrow nfa_accept A x w \rightarrow dfa_accept nfa_to_dfa X w.`

Lemma `nfa_to_dfa_sound` (X : `nfa_to_dfa`) w :

`dfa_accept nfa_to_dfa X w \rightarrow [exists x, (x \in X) && nfa_accept A x w].`

Lemma `nfa_to_dfa_correct` : `nfa_lang A =i dfa_lang nfa_to_dfa .`

4.2 Connected Components

Finite automata can have isolated subsets of states that are not reachable from the starting state. These states can not contribute to the language of the automaton since there are no runs from the starting state to any of those unreachable states. It will later be useful to have automata that only contain reachable states. Therefore, we define a procedure to extract the connected component containing the starting state from a given automaton.

Definition 4.2.1. Let $A = (Q, s, F, \delta)$ be a DFA. We define reachable1 such that for all x and y , $(x, y) \in \text{reachable1} \iff \exists a, (x, a, y) \in \delta$. We define $\text{reachable} := \{y \mid (s, y) \in \text{reachable1}^*\}$, where reachable1^* denotes the transitive closure of reachable1 . With this, we can define the connected automaton A_c :

$$\begin{aligned} Q_c &:= Q \cap \text{reachable} \\ s_c &:= s \\ F_c &:= F \cap \text{reachable} \\ \delta_c &:= \{(x, a, y) \mid (x, a, y) \in \delta \wedge x, y \in Q_c\} \\ A_c &:= (Q_c, s_c, F_c, \delta_c). \end{aligned}$$

We make use of SSREFLECT's *connect* predicate to extract a sequence of all states reachable from s . From this, we construct a finite type and use that as the new set of states. These new states carry a proof of reachability. We also have to give a transition function that ensures transitions always end in reachable states.

Definition $\text{reachable1} := [\text{fun } x \ y \Rightarrow [\text{exists } a, \text{dfa_step } A1 \ x \ a == y]]$.

Definition $\text{reachable} := \text{enum } (\text{connect } \text{reachable1 } (\text{dfa_s } A1))$.

Lemma $\text{reachable0} : \text{dfa_s } A1 \ \backslash \text{in } \text{reachable}$.

Lemma $\text{reachable_step } x \ a : x \ \backslash \text{in } \text{reachable} \rightarrow \text{dfa_step } A1 \ x \ a \ \backslash \text{in } \text{reachable}$.

Definition $\text{dfa_connected} :=$

```
{| dfa_s := SeqSub reachable0;
  dfa_fin := fun x => match x with SeqSub x _ => dfa_fin A1 x end;
  dfa_step := fun x a => match x with
    | SeqSub _ Hx => SeqSub (reachable_step _ a Hx)
  end |}.
```

Make
this
under-
stand-
able

Lemma 4.2.2. For every state $x \in \text{reachable}$ we have that

$$\mathcal{L}_x(A_c) = \mathcal{L}_x(A).$$

Proof. “ \subseteq ” Trivial. “ \supseteq ” We do an induction on w .

- For $w = \varepsilon$ we have $\varepsilon \in \mathcal{L}_x(A)$ and therefore $x \in F$. With $x \in \text{reachable}$ we get $x \in F_c$. Thus, $\varepsilon \in \mathcal{L}_x(A_c)$.

- For $w = aw'$ we have $y \in Q$ such that $(x, a, y) \in \delta$ and $w' \in \mathcal{L}_y(A)$. From $x \in \text{reachable}$ we get $y \in \text{reachable}$ by transitivity. Therefore, $(x, a, y) \in \delta_c$. The induction hypothesis gives us $w' \in \mathcal{L}_y(A_c)$. Thus, $aw' \in \mathcal{L}_x(A_c)$.

□

Theorem 4.2.3. *The language of the connected automaton A_c is identical to that of the original automaton A , i.e.*

$$\mathcal{L}(A) = \mathcal{L}(A_c).$$

Proof. By reflexivity, we have $s \in \text{reachable}$. We use lemma 4.2.2 with $x = s$ to prove the claim. □

The formalization of lemma 4.2.2 and theorem 4.2.3 is straight-forward.

Lemma `dfa_connected_correct' x (Hx: x \in reachable) :`
`dfa_accept dfa_connected (SeqSub Hx) =i dfa_accept A1 x.`

Lemma `dfa_connected_correct: dfa_lang dfa_connected =i dfa_lang A1.`

To make use of the fact that A_c is fully connected, we will prove a characteristic property of A_c .

Definition 4.2.4. A **representative** of a state x is a word w such that the unique run of w on A_c ends in x .

Lemma 4.2.5. *We can give a representative for every state $x \in Q_c$.*

Proof. x carries a proof of reachability. From this, we get a path through the graph of `reachable1` that ends in x . We build the representative by extracting the edges of the path and building a word from those. □

Choice?

The formalization of theorem 4.2.5 includes a more general version of the theorem, which facilitates the proof by induction over the path.

Lemma `dfa_connected_repr' (x y: dfa_connected):`
`connect reachable1_connected y x ->`
`exists w, last y (dfa_run' dfa_connected y w) = x.`

Lemma `dfa_connected_repr x :`
`exists w, last (dfa_s dfa_connected) (dfa_run dfa_connected w) = x.`

4.3 Emptiness

Given an automaton A , we can check if $\mathcal{L}(A) = \emptyset$. We simply obtain the connected automaton of A and check if there are any final states left.

Theorem 4.3.1. *The language of the connected automaton A_c is empty if and only if its set of final states F_c is empty, i.e.*

$$\mathcal{L}(A) = \emptyset \iff F_c = \emptyset.$$

Proof. By theorem 4.2.3 we have $\mathcal{L}(A) = \mathcal{L}(A_c)$. Therefore, it suffices to show

$$\mathcal{L}(A_c) = \emptyset \iff F_c = \emptyset.$$

“ \Leftarrow ” We have $\mathcal{L}(A_c) = \emptyset$ and have to show that for all $x \in Q_c$, $x \notin F_c$. Let $x \in Q_c$. By lemma 4.2.5 we get w such that the unique run of w on A_c ends in x . We use $\mathcal{L}(A_c) = \emptyset$ to get $w \notin \mathcal{L}(A_c)$, which implies that the run of w on A_c ends in a non-final state. By substituting the last state of the run by x we get $x \notin F_c$. “ \Rightarrow ” We have $F_c = \emptyset$ and have to show that for all words w , $w \notin \mathcal{L}(A_c)$. We use $F_c = \emptyset$ to show that the last state of the run of w on A_c is non-final. Thus, $w \notin \mathcal{L}(A_c)$.

Thus, emptiness is decidable. \square

The formalization of lemma 4.3.1 is split in two parts to facilitate its application.

Definition `dfa_lang_empty := #|dfa_fin dfa_connected| == 0.`

Lemma `dfa_lang_empty_complete: dfa_lang dfa_connected =i pred0 -> dfa_lang_empty.`

Lemma `dfa_lang_empty_sound: dfa_lang_empty -> dfa_lang dfa_connected =i pred0.`

Lemma `dfa_lang_empty_correct:`

`reflect (dfa_lang A1 =i pred0)`
`dfa_lang_empty.`

4.4 Deciding Equivalence of Finite Automata

Given finite automata A_1 and A_2 , we construct DFA A such that the language of A is the symmetric difference of the languages of A_1 and A_2 , i.e.,

$$\mathcal{L}(A) := \mathcal{L}(A_1) \ominus \mathcal{L}(A_2) = \mathcal{L}(A_1) \cap \neg \mathcal{L}(A_2) \cup \mathcal{L}(A_2) \cap \neg \mathcal{L}(A_1).$$

Theorem 4.4.1. *The equivalence of A_1 and A_2 is decidable, i.e.*

$$\mathcal{L}(A_1) = \mathcal{L}(A_2) \text{ if and only if } \mathcal{L}(A) \text{ is empty.}$$

Proof. The correctness of this procedure follows from the properties of the symmetric difference operator, i.e.

$$\mathcal{L}(A_1) \ominus \mathcal{L}(A_2) = \emptyset \Leftrightarrow \mathcal{L}(A_1) = \mathcal{L}(A_2).$$

Thus, equivalence is decidable. \square

Listing 4.3: Formalization of theorem 4.4.1

Definition `dfa_sym_diff` :=
`dfa_disj (dfa_conj A1 (dfa_compl A2)) (dfa_conj A2 (dfa_compl A1)).`
Lemma `dfa_sym_diff_correct`:
`dfa_lang_empty dfa_sym_diff <=> dfa_lang A1 =i dfa_lang A2.`

4.5 Regular Expressions and Finite Automata

We prove that there is a finite automaton for every extended regular expression and vice versa. In fact, we can give a standard regular expression for every finite automaton. With this, we will prove that extended regular expressions are equivalent to standard regular expressions, thereby proving closure under intersection and negation.

4.5.1 Regular Expressions to Finite Automata

We prove that there exists an equivalent automaton for every extended regular expression. The structure of this proof is given by the inductive nature of regular expressions. For every constructor, we provide an equivalent automaton.

Depending on the constructor of the regular expression, we will construct an equivalent DFA or NFA. Void, Eps, Dot, Atom, Plus, And and Not are very easy to implement on DFAs, whereas Star and Conc lend themselves well to NFAs.

Void

Definition 4.5.1. *We define an empty DFA with a single, non-accepting state, i.e.*

$$A_{\emptyset} := (\{t\}, t, \emptyset, \{(t, a, t) \mid a \in \Sigma\}).$$

Lemma 4.5.2. *The language of the empty DFA is empty, i.e.*

$$\mathcal{L}(E) = \emptyset.$$

Proof. A_{\emptyset} has no final states, i.e. no run can end in a final state. □

Definition `dfa_void` :=
`{| dfa_s := tt;`
`dfa_fin := pred0;`
`dfa_step := [fun x a => tt] |}.`

Lemma `dfa_void_correct` $x w$: $\sim\sim$ `dfa_accept dfa_void` $x w$.

Eps

Definition 4.5.3. We define an automaton that accepts only the empty word, i.e.

$$A_\varepsilon := (\{t, f\}, t, \{f\}, \{(x, a, f) \mid x \in \{t, f\}, a \in \Sigma\}).$$

Lemma 4.5.4. A_ε accepts no word in state f , i.e. for all w ,

$$w \notin \mathcal{L}_f(A_\varepsilon).$$

Proof. Let $w \in \Sigma^*$. We do an induction on w . For $w = \varepsilon$ we get $\varepsilon \notin \mathcal{L}_f(A_\varepsilon)$ by $f \notin F_\varepsilon$. For $w = aw'$ we have $w' \notin \mathcal{L}_f(A_\varepsilon)$. Furthermore, $(f, a, f) \in \delta_\varepsilon$. Therefore, $aw' \notin \mathcal{L}_f(A_\varepsilon)$. \square

Lemma 4.5.5. The language of A_ε is exactly the singleton set containing the empty word, i.e.

$$\mathcal{L}(A_\varepsilon) = \{\varepsilon\}.$$

Proof. Let $w \in \Sigma^*$. We do an induction on w . For $w = \varepsilon$ we have $\varepsilon \in \mathcal{L}(A_\varepsilon)$ and $\varepsilon \in \{\varepsilon\}$. Therefore, both directions are trivial. For $w = aw'$ we consider both directions independently.

“ \Rightarrow ” We have $(t, a, f) \in \delta_\varepsilon$ and $w' \in \mathcal{L}_f(A_\varepsilon)$. By lemma 4.5.4, this is a contradiction.

“ \Leftarrow ” We get a straight-forward contradiction from $aw' \in \{\varepsilon\}$. \square

Definition `dfa_eps :=`

```
{ | dfa_s := true;
    dfa_fin := pred1 true;
    dfa_step := [fun x a => false] |}.
```

Lemma `dfa_eps.correct: dfa_lang dfa_eps =i pred1 [::].`

reflect?

Dot

Definition 4.5.6. We define an automaton that accepts the set of all singleton words, i.e.

$$A_{Dot} := (\{s, t, f\}, s, \{t\}, \{(s, a, t) \mid a \in \Sigma\} \cup \{(x, a, f) \mid x \in \{t, f\}, a \in \Sigma\}).$$

Lemma 4.5.7. A_{Dot} does not accept any word in state f , i.e. $\mathcal{L}_f(A_{Dot}) = \emptyset$.

Proof. We prove this by induction on $w \in \Sigma^*$. For $w = \varepsilon$ we have $\varepsilon \notin \mathcal{L}_f(A_{Dot})$ by $f \notin F_{Dot}$. For $w = aw'$ we have $(f, a, f) \in \delta_{Dot}$ and $w' \notin \mathcal{L}_f(A_{Dot})$ by induction hypothesis. Thus, $aw' \notin \mathcal{L}_f(A_{Dot})$. \square

Lemma 4.5.8. A_{Dot} accepts exactly the empty word in state t , i.e. $\mathcal{L}_t(A_{Dot}) = \{\varepsilon\}$.

Proof. Let $w \in \Sigma^*$. We do a case distinction on w . For $w = \varepsilon$ we have $\varepsilon \in \mathcal{L}_t(A_{Dot})$ by $t \in F_{Dot}$. We also have $\varepsilon \in \{\varepsilon\}$. For $w = aw'$ we get $(t, a, f) \in \delta_{Dot}$. Since $aw' \notin \{\varepsilon\}$ it suffices to show that $w' \notin \mathcal{L}_f(A_{Dot})$, which we have by lemma 4.5.7. \square

Definition `dfa_dot` :=

```
{| dfa_s := None;
   dfa_fin := pred1 (Some true);
   dfa_step := [fun x b => if x == None then Some true else Some false ] |}.
```

Lemma `dfa_dot_correct''` w: $\sim\sim$ `dfa_accept dfa_dot` (Some false) w.

Lemma `dfa_dot_correct'` w: `dfa_accept dfa_dot` (Some true) w = (w == [:]).

Lemma `dfa_dot_correct` w: `dfa_lang dfa_dot` w = (size w == 1).

reflect?

Not

Definition 4.5.9. Given DFA $A = (Q, s, F, \delta)$, the complement automaton A_{\neg} is constructed by switching accepting and non-accepting states, i.e.

$$A_{\neg} := (Q, s, Q \setminus F, \delta).$$

Lemma 4.5.10. For every state $x \in Q$, we have that $w \in \Sigma^*$ is accepted in x by A_{\neg} if and only if it is not accepted in x by A , i.e. $\mathcal{L}_x(A_{\neg}) = \Sigma^* \setminus \mathcal{L}_x(A)$

Proof. We do an induction on w . For $w = \varepsilon$ we have $\varepsilon \in \mathcal{L}_x(A_{\neg}) \iff \varepsilon \in \mathcal{L}_x(A)$ from $x \in F \iff x \notin Q \setminus F$. For $w = aw'$ we get $(y, a, x) \in \delta$. By induction hypothesis, $w' \in \mathcal{L}_x(A_{\neg}) \iff w' \notin \mathcal{L}_x(A)$. Thus, $aw' \in \mathcal{L}_y(A_{\neg}) \iff aw' \notin \mathcal{L}_y(A)$. \square

Lemma 4.5.11. A_{\neg} accepts the complement language of A , i.e. $\mathcal{L}(A_{\neg}) = \Sigma^* \setminus \mathcal{L}(A)$.

Proof. This follows directly from lemma 4.5.10 with $x = s$. \square

Plus

Definition 4.5.12. Given DFAs $A_1 = (Q_1, s_1, F_1, \delta_1)$ and $A_2 = (Q_2, s_2, F_2, \delta_2)$ we construct the disjunction automaton in the following way:

$$\begin{aligned} Q_{\vee} &:= Q_1 \times Q_2 \\ s_{\vee} &:= (s_1, s_2) \\ F_{\vee} &:= \{(x_1, x_2) \mid x_1 \in F_1 \vee x_2 \in F_2\} \\ \delta_{\vee} &:= \{((x_1, x_2), a, (y_1, y_2)) \mid a \in \Sigma, (x_1, a, y_1) \in \delta_1, (x_2, a, y_2) \in \delta_2\} \\ A_{\vee} &:= (Q_{\vee}, s_{\vee}, F_{\vee}, \delta_{\vee}). \end{aligned}$$

Lemma 4.5.13. *For every state $(x_1, x_2) \in Q_\vee$, we have that*

$$\mathcal{L}_{(x_1, x_2)}(A_\vee) = \mathcal{L}_{x_1}(A_1) \cup \mathcal{L}_{x_2}(A_2).$$

Proof. We do a proof by induction on $w \in \Sigma^*$. For $w = \varepsilon$ we have, by definition of F_\vee , $\varepsilon \in \mathcal{L}_{(x_1, x_2)}(A_\vee) \iff \varepsilon \in \mathcal{L}_{x_1}(A_1) \vee \varepsilon \in \mathcal{L}_{x_2}(A_2)$. For $w = aw'$ we get $(x_1, a, y_1) \in \delta_1$ and $(x_2, a, y_2) \in \delta_2$. By induction hypothesis, we also have $w' \in \mathcal{L}_{x_1}(A_1) \vee w' \in \mathcal{L}_{x_2}(A_2)$. Thus, we get $aw' \in \mathcal{L}_{y_1}(A_1) \vee aw' \in \mathcal{L}_{y_2}(A_2)$. \square

Lemma 4.5.14. $\mathcal{L}(A_\vee) = \mathcal{L}(A_1) \cup \mathcal{L}(A_2)$.

Proof. This follows directly from lemma 4.5.13 with $x = (s_1, s_2)$. \square

Definition `dfa_disj` :=

```
{| dfa_s := (dfa_s A1, dfa_s A2);
  dfa_fin := (fun q => let (x1,x2) := q in dfa_fin A1 x1 || dfa_fin A2 x2);
  dfa_step := [fun x a => (dfa_step A1 x.1 a, dfa_step A2 x.2 a)] |}.
```

Lemma `dfa_disj_correct'` x:

```
[ predU dfa_accept A1 x.1 & dfa_accept A2 x.2 ]
=i dfa_accept dfa_disj x.
```

Lemma `dfa_disj_correct`:

```
[ predU dfa_lang A1 & dfa_lang A2 ]
=i dfa_lang dfa_disj .
```

And

Definition 4.5.15. *Given DFAs $A_1 = (Q_1, s_1, F_1, \delta_1)$ and $A_2 = (Q_2, s_2, F_2, \delta_2)$ we construct the conjunction automaton in the following way:*

$$\begin{aligned} Q_\wedge &:= Q_1 \times Q_2 \\ s_\wedge &:= (s_1, s_2) \\ F_\wedge &:= \{(x_1, x_2) \mid x_1 \in F_1 \wedge x_2 \in F_2\} \\ \delta_\wedge &:= \{((x_1, x_2), a, (y_1, y_2)) \mid a \in \Sigma, (x_1, a, y_1) \in \delta_1, (x_2, a, y_2) \in \delta_2\} \\ A_\wedge &:= (Q_\wedge, s_\wedge, F_\wedge, \delta_\wedge). \end{aligned}$$

Lemma 4.5.16. *For every state $(x_1, x_2) \in Q_\wedge$, we have that*

$$\mathcal{L}_{(x_1, x_2)}(A_\wedge) = \mathcal{L}_{x_1}(A_1) \cap \mathcal{L}_{x_2}(A_2).$$

Proof. This proof is very similar to lemma 4.5.13. We do a proof by induction on $w \in \Sigma^*$. For $w = \varepsilon$ we have, by definition of F_\wedge , $\varepsilon \in \mathcal{L}_{(x_1, x_2)}(A_\wedge) \iff \varepsilon \in \mathcal{L}_{x_1}(A_1) \wedge \varepsilon \in \mathcal{L}_{x_2}(A_2)$. For $w = aw'$ we get $(x_1, a, y_1) \in \delta_1$ and $(x_2, a, y_2) \in \delta_2$. By induction hypothesis, we also have $w' \in \mathcal{L}_{x_1}(A_1) \wedge w' \in \mathcal{L}_{x_2}(A_2)$. Thus, we get $aw' \in \mathcal{L}_{y_1}(A_1) \wedge aw' \in \mathcal{L}_{y_2}(A_2)$. \square

Lemma 4.5.17. $\mathcal{L}(A_\wedge) = \mathcal{L}(A_1) \cap \mathcal{L}(A_2)$.

Proof. This follows directly from lemma 4.5.16 with $x = (s_1, s_2)$. \square

Definition `dfa_conj` :=

```
{| dfa_s := (dfa_s A1, dfa_s A2);
  dfa_fin := (fun x => dfa_fin A1 x.1 && dfa_fin A2 x.2);
  dfa_step := [fun x a => (dfa_step A1 x.1 a, dfa_step A2 x.2 a)] |}.
```

Lemma `dfa_conj_correct' x1 x2` :

```
[ predI dfa_accept A1 x1 & dfa_accept A2 x2 ]
=i dfa_accept dfa_conj (x1, x2).
```

Lemma `dfa_conj_correct`:

```
[ predI dfa_lang A1 & dfa_lang A2 ]
=i dfa_lang dfa_conj.
```

Conc

Definition 4.5.18. Given two NFAs $A_1 = (Q_1, s_1, F_1, \delta_1)$ and $A_2 = (Q_2, s_2, F_2, \delta_2)$ we construct the concatenation automaton in the following way:

$$\begin{aligned} Q_{Conc} &:= Q_1 \cup Q_2 \\ s_{Conc} &:= s_1 \\ F_{Conc} &:= \text{TODO} \\ \delta_{Conc} &:= \delta_1 \cup \delta_2 \cup \{(x, a, y) \mid x \in Q_1, y \in Q_2, (s_2, a, y) \in \delta_2\} \\ A_{Conc} &:= (Q_{Conc}, s_{Conc}, F_{Conc}, \delta_{Conc}). \end{aligned}$$

Case
dist. on
 $s_2 \in F_2$

Lemma 4.5.19. Every run of A_2 can be mapped to a run in A_3 .

Proof. Let σ be a run starting in x for $w \in \Sigma^*$ on A_2 . We do an induction on σ . For $|\sigma| = 0$ we have $w = \varepsilon$. Therefore, we have that σ is also a run starting in x for ε on A_{Conc} . For $\sigma = y\sigma'$ we have $w = aw'$, $(x, a, y) \in \delta_2$. By definition of δ_{Conc} we also have $(x, a, y) \in \delta_{Conc}$. By induction hypothesis, we have that σ' is a run for w' starting in y on A_{Conc} . Thus, $y\sigma'$ is a run for aw' starting in x on A_{Conc} . \square

symbol
for
empty
run

Include
all
proofs

4.5.2 Deciding Equivalence of Regular Expressions

Based on our procedure to construct an equivalent automaton from a regular expression, we can decide equivalence of regular expressions. Given r_1 and r_2 , we construct equivalent DFA A_1 and A_2 as above.

4.5.3 Finite Automata to Regular Expressions

We prove that there is an equivalent standard regular expression for every finite automaton. There are three ways to prove this.

The first one is a method called “state removal” [11] (reformulated in [17]), which works by sequentially building up regular expressions on the edges between states. In every step, one state is removed and its adjacent states’ edges are updated to incorporate the missing state into the regular expression. Finally, only two states remain. The resulting edges can be combined to form a regular expression that recognizes the language of the initial automaton.

The second method is known as “Brzozowski’s method” [12] and builds upon Brzozowski derivatives of regular expressions. This method is algebraic in nature and arrives at a regular expression by solving a system of linear equations on regular expressions. Every state is assigned an unknown regular expression. The intuition of these unknown regular expressions is that they recognize the words accepted in their associated state. The system is solved by substitution and Arden’s lemma [3]. The regular expressions associated with the starting state recognizes the language of the automaton.

The third method, which we chose for our development, is due to Kleene [22]. It is known as the “transitive closure method”. This method recursively constructs a regular expression that is equivalent to the given automaton. For the remainder of the chapter, we assume that we are given a DFA $A = (Q, s, F, \delta)$.

The idea of the “transitive closure method” is that we can give a regular expression to describe the path between any two states x and y . This regular expression accepts all words whose run σ on A starting in x ends in y . In fact, we can even give such a regular expression if we limit the set of paths through which the run is allowed to pass. We will call this set X . Passing through, here, means that the restriction applies only to states that are traversed, i.e. not to the beginning or end of the run (x and y in this case).

If we take X to be the empty set, we only consider two types of runs. First, if $x \neq y$, every transition from x to y constitutes one (singleton) word. Conversely, if there is a word which does not pass through a state and whose run on A starts in x and ends in y , it can only be a singleton word consisting of one of the transitions from x to y . Therefore, the corresponding regular expression is the disjunction of all transitions from x to y . These transitions constitutes all possible words that lead from x to y without passing through any state.

Otherwise, if $x = y$, we also have to consider the empty word, since its run on A starts in x and ends in y . Thus, the corresponding regular expression is the disjunction of all transitions from x to y **and** ε .

In the case of a non-empty X , we make the following observation. If we pick an element $z \in X$, then any run σ from x to y either passes through z , or does not pass through z . If it does, we can split it into three parts.

- (i) The first part contains the prefix of σ which contains all states up to the first occurrence of z that is not the starting state.
- (ii) The second part contains that part of the remainder of σ which contains all further occurrences of z , though not the last state if that is z .
- (iii) The third part contains the remainder.

Parts (i) and (iii) can easily be expressed in terms of $X \setminus \{z\}$. Part (ii) can be further decomposed into runs from z to z that do not pass through z . Thus, part (ii) can also be expressed in terms of $X \setminus \{z\}$ with the help of the $*$ operator.

If σ does not pass through z , it is covered by the regular expression for paths from x to y restricted to $X \setminus \{z\}$.

Definition 4.5.20. Let $X \subseteq Q$. Let $x, y \in Q$. We define R recursively on $|X|$:

correct?

$$R_{x,y}^X := \begin{cases} \sum_{\substack{a \in \Sigma \\ (x,a,y) \in \delta}} a & \text{if } X = \emptyset \wedge x \neq y; \\ \sum_{\substack{a \in \Sigma \\ (x,a,y) \in \delta}} a + \varepsilon & \text{if } X = \emptyset \wedge x = y; \\ R_{x,z}^{X \setminus \{z\}} (R_{z,z}^{X \setminus \{z\}})^* R_{z,y}^{X \setminus \{z\}} + R_{x,y}^{X \setminus \{z\}} & \text{if } \exists z \in X. \end{cases}$$

Based on the intuition given above, we try to express $\mathcal{L}(A)$ by R . The goal of this chapter is to prove the following theorem.